

Rate Limiting Algorithms Explained with Code

#20 Rate Limiting Algorithms



ASHISH PRATAP SINGH

JUL 17, 2024




77



Share


Rate limiting helps protect services from being overwhelmed by too many requests from a single user or client.



In this article we will dive into 5 of the most common rate limiting algorithms, their pros and cons and learn how to implement them in code.

Rate Limiting Algorithms

1. Token Bucket



The Token Bucket algorithm is one of the most popular and widely used rate limiting approaches due to its simplicity and effectiveness.

How It Works:

- Imagine a bucket that holds tokens.
- The bucket has a maximum capacity of tokens.
- Tokens are added to the bucket at a fixed rate (e.g., 10 tokens per second).
- When a request arrives, it must obtain a token from the bucket to proceed.
- If there are enough tokens, the request is allowed and tokens are removed.
- If there aren't enough tokens, the request is dropped.

```
import time

class TokenBucket:
    def __init__(self, capacity, fill_rate):
        self.capacity = capacity # Maximum number of tokens the bucket can hold
        self.fill_rate = fill_rate # Rate at which tokens are added (tokens/second)
        self.tokens = capacity # Current token count, start with a full bucket
        self.last_time = time.time() # Last time we checked the token count

    def allow_request(self, tokens=1):
        now = time.time()
        # Calculate how many tokens have been added since the last check
        time_passed = now - self.last_time
        self.tokens = min(self.capacity, self.tokens + time_passed * self.fill_rate)
        self.last_time = now

        # Check if we have enough tokens for this request
        if self.tokens >= tokens:
            self.tokens -= tokens
            return True
        return False

# Usage example
limiter = TokenBucket(capacity=10, fill_rate=1) # 10 tokens, refill 1 token per second

for _ in range(15):
    print(limiter.allow_request()) # Will print True for the first 10 requests, then False
    time.sleep(0.1) # Wait a bit between requests

time.sleep(5) # Wait for bucket to refill
print(limiter.allow_request()) # True
```

Code Link: [Python](#), [Java](#)

Pros:

- Relatively straightforward to implement and understand.
- Allows bursts of requests up to the bucket's capacity, accommodating short-term spikes.

Cons:

- The memory usage scales with the number of users if implemented per-user.
- It doesn't guarantee a perfectly smooth rate of requests.



If you want to take your **Software Engineering** skills to the next level, check out [CodeCrafters](#).

CodeCrafters is a YC backed platform that offers a unique, **hands-on** approach to practice **complex programming challenges** like building your own **Redis**, **Git**, **SQLite** and **Bittorrent** client from scratch.

[**Sign up and get 40% off if you upgrade.**](#)

2. Leaky Bucket



The Leaky Bucket algorithm is similar to Token Bucket but focuses on smoothing out bursty traffic.

How it works:

1. Imagine a bucket with a small hole in the bottom.
2. Requests enter the bucket from the top.
3. The bucket processes ("leaks") requests at a constant rate through the hole.
4. If the bucket is full, new requests are discarded.

```
from collections import deque
import time

class LeakyBucket:
    def __init__(self, capacity, leak_rate):
        self.capacity = capacity # Maximum number of requests in the bucket
        self.leak_rate = leak_rate # Rate at which requests leak (requests/second)
        self.bucket = deque() # Queue to hold request timestamps
        self.last_leak = time.time() # Last time we leaked from the bucket

    def allow_request(self):
        now = time.time()
        # Simulate leaking from the bucket
        leak_time = now - self.last_leak
        leaked = int(leak_time * self.leak_rate)
        if leaked > 0:
            # Remove the leaked requests from the bucket
            for _ in range(min(leaked, len(self.bucket))):
                self.bucket.popleft()
            self.last_leak = now

        # Check if there's capacity and add the new request
        if len(self.bucket) < self.capacity:
            self.bucket.append(now)
            return True
        return False

# Usage example
limiter = LeakyBucket(capacity=5, leak_rate=1) # 5 requests, leak 1 per second

for _ in range(10):
    print(limiter.allow_request()) # Will print True for the first 5 requests, then False
    time.sleep(0.1) # Wait a bit between requests

time.sleep(1) # Wait for bucket to leak
```

```
print(limiter.allow_request()) # True
```

Code Link: [Python](#), [Java](#)


Pros:

- Processes requests at a steady rate, preventing sudden bursts from overwhelming the system.
- Provides a consistent and predictable rate of processing requests.

Cons:

- Does not handle sudden bursts of requests well; excess requests are immediately dropped.
- Slightly more complex to implement compared to Token Bucket.

3. Fixed Window Counter



The Fixed Window Counter algorithm divides time into fixed windows and counts requests in each window.

How it works:

1. Time is divided into fixed windows (e.g., 1-minute intervals).
2. Each window has a counter that starts at zero.
3. New requests increment the counter for the current window.
4. If the counter exceeds the limit, requests are denied until the next window.

```
import time

class FixedWindowCounter:
    def __init__(self, window_size, max_requests):
        self.window_size = window_size # Size of the window in seconds
        self.max_requests = max_requests # Maximum number of requests per window
        self.current_window = time.time() // window_size
        self.request_count = 0

    def allow_request(self):
        current_time = time.time()
        window = current_time // self.window_size

        # If we've moved to a new window, reset the counter
        if window != self.current_window:
            self.current_window = window
            self.request_count = 0

        # Check if we're still within the limit for this window
        if self.request_count < self.max_requests:
            self.request_count += 1
            return True
        return False

# Usage example
limiter = FixedWindowCounter(window_size=60, max_requests=5) # 5 requests per minute

for _ in range(10):
    print(limiter.allow_request()) # Will print True for the first 5 requests, then False
    time.sleep(0.1) # Wait a bit between requests

time.sleep(60) # Wait for the window to reset
print(limiter.allow_request()) # True
```

Code Link: [Python](#), [Java](#)

Pros:


- Easy to implement and understand.
- Provides clear and easy-to-understand rate limits for each time window.

Cons:

- Does not handle bursts of requests at the boundary of windows well. Can allow twice the rate of requests at the edges of windows.

Subscribe to receive new articles every week.

4. Sliding Window Log



The Sliding Window Log algorithm keeps a log of timestamps for each request and uses this to determine if a new request should be allowed.

How it works:

1. Keep a log of request timestamps.
2. When a new request comes in, remove all entries older than the window size.
3. Count the remaining entries.
4. If the count is less than the limit, allow the request and add its timestamp to the log.
5. If the count exceeds the limit, request is denied.

```
import time
from collections import deque

class SlidingWindowLog:
    def __init__(self, window_size, max_requests):
        self.window_size = window_size # Size of the sliding window in seconds
        self.max_requests = max_requests # Maximum number of requests per window
        self.request_log = deque() # Log to keep track of request timestamps

    def allow_request(self):
        now = time.time()

        # Remove timestamps that are outside the current window
        while self.request_log and now - self.request_log[0] >= self.window_size:
            self.request_log.popleft()

        # Check if we're still within the limit
        if len(self.request_log) < self.max_requests:
            self.request_log.append(now)
            return True
        return False

# Usage example
limiter = SlidingWindowLog(window_size=60, max_requests=5) # 5 requests per minute

for _ in range(10):
    print(limiter.allow_request()) # Will print True for the first 5 requests, then False
    time.sleep(0.1) # Wait a bit between requests

time.sleep(60) # Wait for the window to slide
print(limiter.allow_request()) # True
```

Code Link: [Python](#), [Java](#)


Pros:

- Very accurate, no rough edges between windows.
- Works well for low-volume APIs.

Cons:

- Can be memory-intensive for high-volume APIs.
- Requires storing and searching through timestamps.

5. Sliding Window Counter



This algorithm combines the Fixed Window Counter and Sliding Window Log approaches for a more accurate and efficient solution.

Instead of keeping track of every single request's timestamp as the sliding log does, it focus on the number of requests from the last window.

So, if you are in 75% of the current window, 25% of the weight would come from the previous window, and the rest from the current one:

```
weight = (100 - 75)% * lastWindowRequests + currentWindowRequests
```

Now, when a new request comes, you add one to that weight (weight + 1). If this new total crosses our set limit, we have to reject the request.

How it works:

1. Keep track of request count for the current and previous window.
2. Calculate the weighted sum of requests based on the overlap with the sliding window.
3. If the weighted sum is less than the limit, allow the request.

```

import time

class SlidingWindowCounter:
    def __init__(self, window_size, max_requests):
        self.window_size = window_size # Size of the sliding window in seconds
        self.max_requests = max_requests # Maximum number of requests per window
        self.current_window = time.time() // window_size
        self.request_count = 0
        self.previous_count = 0

    def allow_request(self):
        now = time.time()
        window = now // self.window_size

        # If we've moved to a new window, update the counts
        if window != self.current_window:
            self.previous_count = self.request_count
            self.request_count = 0
            self.current_window = window

        # Calculate the weighted request count
        window_elapsed = (now % self.window_size) / self.window_size
        threshold = self.previous_count * (1 - window_elapsed) + self.request_count

        # Check if we're within the limit
        if threshold < self.max_requests:
            self.request_count += 1
            return True
        return False

# Usage example
limiter = SlidingWindowCounter(window_size=60, max_requests=5) # 5 requests per minute

for _ in range(10):
    print(limiter.allow_request()) # Will print True for the first 5 requests, then gradually
    become False

```



```
time.sleep(0.1) # wait a bit between requests

time.sleep(30) # Wait for half the window to pass
print(limiter.allow_request()) # Might be True or False depending on the exact timing
```

Code Link: [Python](#), [Java](#)

Pros:

- More accurate than Fixed Window Counter.
- More memory-efficient than Sliding Window Log.
- Smooths out edges between windows.

Cons:

- Slightly more complex to implement.

When implementing rate limiting, consider factors such as the scale of your system, the nature of your traffic patterns, and the granularity of control you need.

Lastly, always communicate your rate limits clearly to your API users, preferably through response headers, so they can implement appropriate retry and backoff strategies in their clients.

Thank you so much for reading.

If you found it valuable, hit a like ❤️ and consider subscribing for more such content every week.

If you have any questions or suggestions, leave a comment.

This post is public so feel free to share it.

Subscribe for free to receive new articles every week.

Checkout my [Youtube channel](#) for more in-depth content.

Follow me on [LinkedIn](#) and [X](#) to stay updated.

Checkout my [GitHub repositories](#) for free interview preparation resources.

I hope you have a lovely day!

See you soon,

Ashish



77 Likes · 8 Restacks

← Previous

Next →

Discussion about this post

Comments

Restacks



Write a comment...



© 2024 Ashish Pratap Singh · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture